

LSTM and its variants




PUNEET KUMAR JAIN, PhD

National Institute of Technology Rourkela


References

- Daniel Jurafsky, James H. Martin, “Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition”, 2nd Ed., PEARSON, 2013
- Several slides adapted from: Chetan Arora, IIT Delhi; Vineeth Balasubramanian, IIT Hyderabad; and others
- Basics of Data Science: <https://learncloudbits.com/post/5-core-steps-to-understand-machine-learning-workflow>


Feed-Forward NN vs LSTM



Feed-forward NN




Simple recurrent networks



long short-term memory NN

Challenges with RNN: Long-Term Dependencies


- The hidden layer in SRNs are being asked to perform two tasks simultaneously:
 - Provide information useful to the decision being made in the current context, and
 - Updating and carrying forward information useful for future decisions.
- A second difficulty (**vanishing gradients**) to successfully training simple recurrent networks arises during the backward pass of training



- **LSTM**

- Selective Read
- Selective write
- Selective Forget

Selective read, write, and forget



- The state (s_i) of an RNN records information from all previous time steps
- At each new timestep the old information gets morphed by the current input
- One could imagine that after t steps the information stored at time step $t - k$ (for some $k < t$) gets completely morphed

so much that it would be impossible to extract the original information stored at time step $t - k$

Selective read, write, and forget




- Let us see an analogy for this
- We can think of the state as a fixed size memory
- Compare this to a fixed size white board that you use to record information
- At each time step (periodic intervals) we keep writing something to the board
- Effectively at each time step we morph the information recorded till that time point
- After many timesteps it would be impossible to see how the information at time step $t - k$ contributed to the state at timestep t

Selective read, write, and forget



- Continuing our whiteboard analogy, suppose we are interested in deriving an expression on the whiteboard
- We follow the following strategy at each time step
- Selectively write on the board
- Selectively read the already written content
- Selectively forget (erase) some content
- Let us look at each of these in detail


LSTM Networks



LSTM Operations: Forget

- First step is to decide what information to throw away from the cell state.
- A sigmoid layer names “forget gate layer” makes this decision.
- It looks at past state output, h_{t-1} and current input, x_t and outputs a number between 0(Forget) and 1(Keep) telling how much to keep.

$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$




LSTM Operations: Input/Insert

- Next step is to decide what new information should be stored in the cell state.
- First, a sigmoid layer called the “input gate layer” decides which values should be updated, i_t . 0(not important) 1(important)
- Next, a ***tanh*** layer creates a vector of new candidate values, \tilde{c}_t that could be added to the state

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$


$$\tilde{c}_t = \tanh(W_c \cdot [h_{t-1}, x_t] + b_c)$$



LSTM Operations: Update

- Old cell state, c_{t-1} is updated into the new cell state, c_t .
- Old state is multiplied by forget layer output, f_t .
- Input gate layer output, it is multiplied with candidate values, \tilde{c}_t and the result is added to values obtained by above multiplication.
- The output of above computations is the new candidate value, c_t .


$$c_t = f_t * c_{t-1} + i_t * \tilde{c}_t$$



LSTM Operations: Output/Recurrent


- **Final step is to decide what to output.**
- Output is based on the current cell state, c_t .
- First a sigmoid layer decides what parts of hidden state is going to output.
- Then cell state is passed through a ***tanh*** layer.
- The output is then multiplied by the output of the sigmoid gate.

$$h_t = o_t * \tanh(c_t)$$

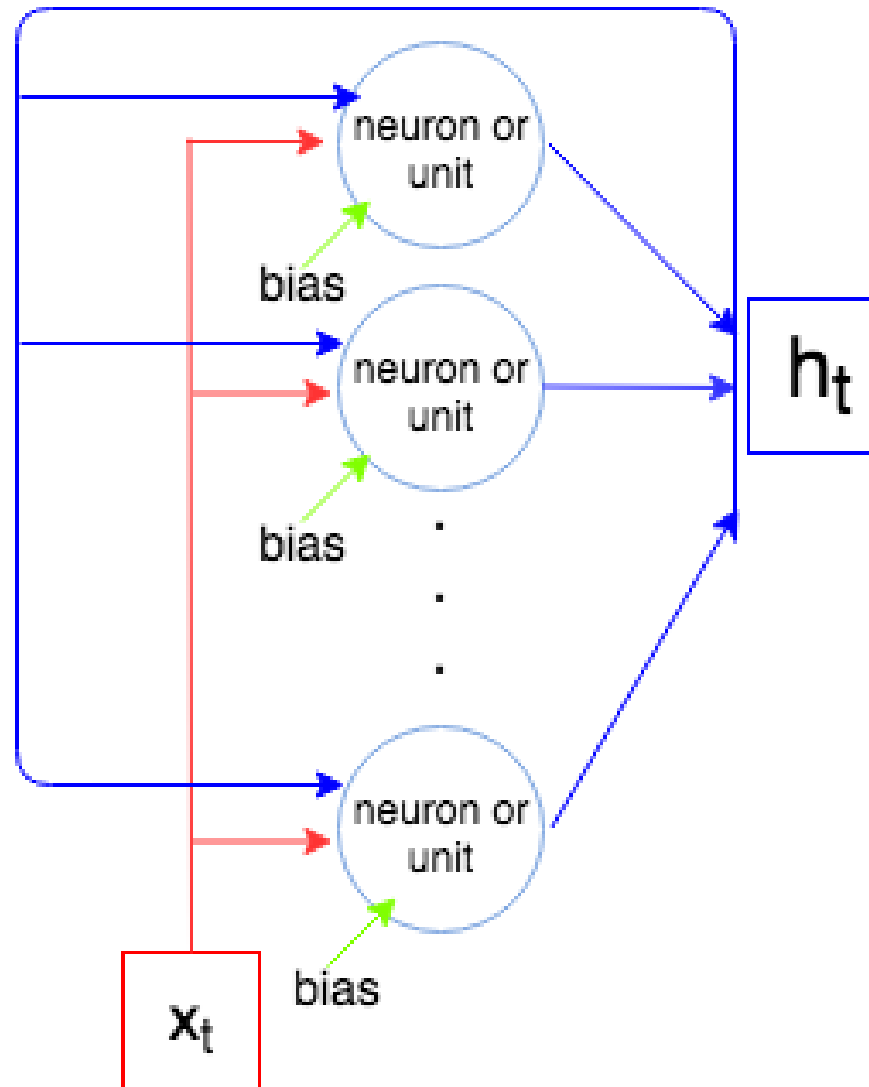


LSTM Architecture

- Each line carries an entire vector.
- Ovals represent pointwise operations, like vector addition.
- Solid rectangles are learned neural network layers.
- Lines merging denote concatenation of vectors.
- Lines splitting denotes vectors being copied and copies going to different locations.



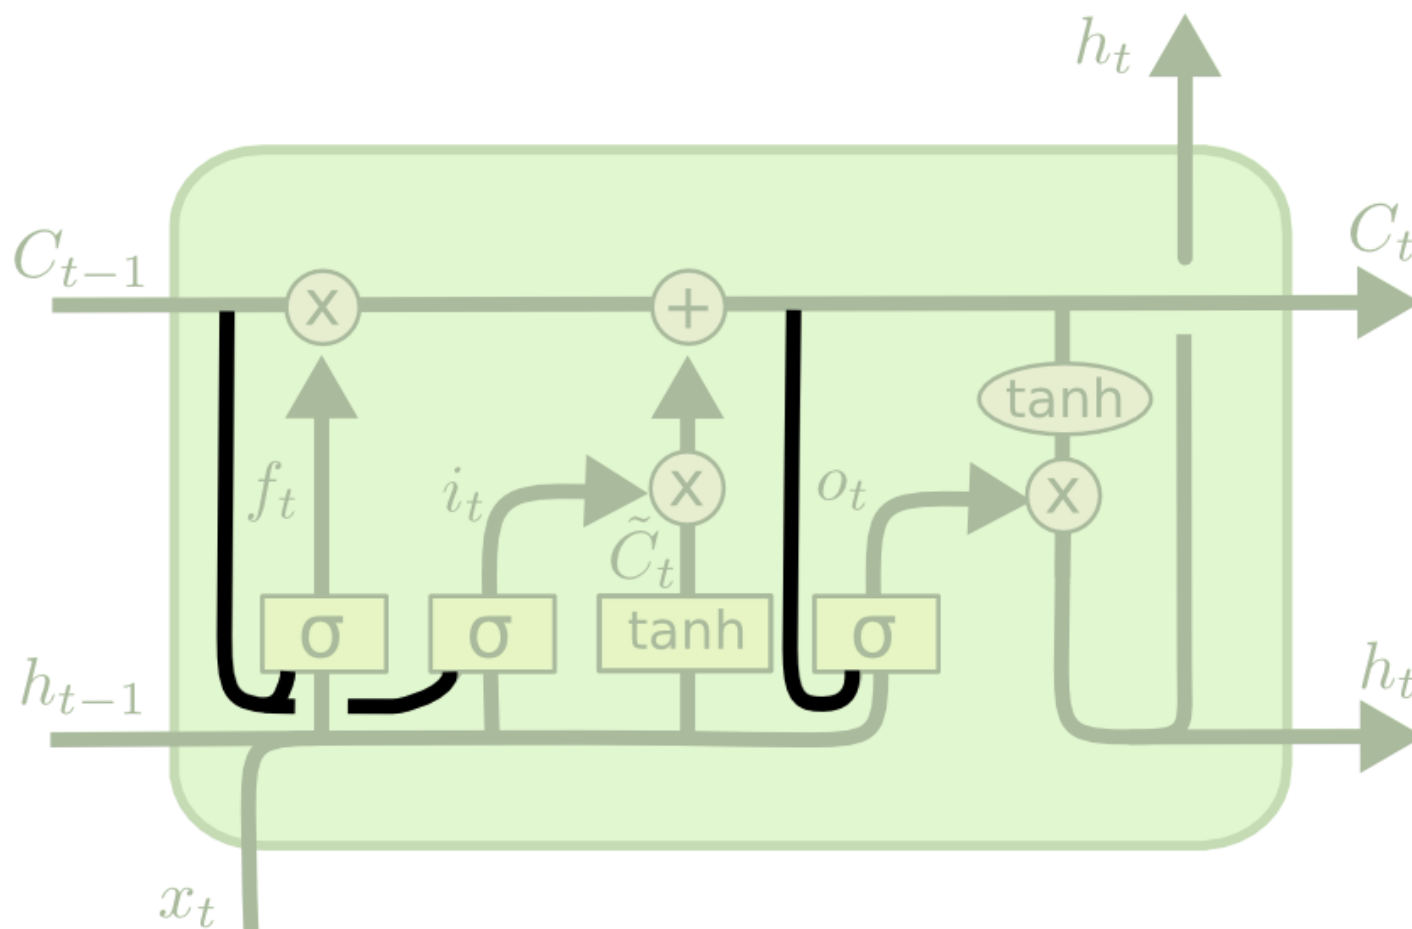
Number of units in hidden layer



Variations of LSTM Architectures


■ The Peephole Variation

- allows the gate layers to read data from the cell state.
- you could also add peepholes to some gates and not other gates.



Variations of LSTM Architectures



- **The Coupled Gate Variation**
- the model makes the decision of what to forget and what to add new information to together



Variations of LSTM Architectures

■ Gated Recurrent Unit (GRU)


- The GRU is like a long short-term memory (LSTM) with a forget gate, but has fewer parameters than LSTM, as it lacks an output gate.
- GRU's performance on certain tasks of polyphonic music modeling, speech signal modeling and natural language processing was found to be similar to that of LSTM.
- GRUs have been shown to exhibit better performance on certain smaller and less frequent datasets.



https://en.wikipedia.org/wiki/Gated_recurrent_unit

Variations of LSTM Architectures

- Bidirectional long short term memory (bi-lstm)
 - processes the data in both forward and backward direction.



Conclusions

- Machine Learning for Data Science plays vital role in daily life
- LSTM are essential NN architecture for sequential data
- LSTM and GRU and derivatives are able to learn a lot of longer term information!
- RNN are not hardware friendly
- LSTMs are **prone to overfitting** and it is difficult to apply the dropout algorithm to curb this issue.



- Thanks

- Puneet Kumar Jain
jainp@nitrkl.ac.in